

Evaluation of performance and productivity metrics of potential programming languages in the HPC environment

— Bachelor Thesis —

Division Scientific Computing
Department of Informatics
Faculty of Mathematics, Informatics und Natural Sciences
University of Hamburg

Submitted by:	Florian Wilkens
E-Mail:	1wilkens@informatik.uni-hamburg.de
Matriculation number:	6324030
Course of studies:	Software-System-Entwicklung
First assessor:	Prof. Dr. Thomas Ludwig
Second assessor:	Sandra Schröder
Advisor:	Michael Kuhn, Sandra Schröder

Hamburg, March 19, 2015

Abstract

This thesis aims to analyze new programming languages in the context of HPC. To compare not only speed but also development productivity and general inner metrics, a basic traffic simulation is implemented in C, Mozilla's Rust and Google's Go. These two languages were chosen on their basic promise of performance as well as memory-safety in the case of Rust or easy multithreaded execution (Go). The implementations are limited to shared-memory parallelism to achieve a fair comparison since the library support for inter-process communication is rather limited at the moment.

Nonetheless the comparison should allow a decent rating of the viability of these two languages in high-performance computing.

Table of Contents

1	Introduction	4
1.1	Motivation	4
1.2	Goals of this thesis	5
2	State of the art	6
2.1	Programming paradigms in Fortran and C	6
2.2	Language candidates	6
2.3	Related work	13
3	Approach	14
3.1	Overview: Streets4MPI	14
3.2	Differences to the base implementation	14
3.3	Implementation process	16
4	Implementation	18
4.0	Project setup	18
4.1	Counting nodes, ways and relations in an .osm.pbf file	21
4.2	Building a basic graph representation for the simulation	21
5	Evaluation	23
6	Conclusion	24
	Bibliography	25
	List of Figures	27
	List of Tables	28
	List of Listings	29
	Appendices	30
A	System configuration	31

1. Introduction

This chapter provides some background information to high-performance computing. The first section describes problems with the currently used programming languages and motivates the search for new candidates. After that the chapter concludes with a quick rundown of the thesis' goals.

1.1. Motivation

The world of high-performance computing is evolving rapidly and programming languages used in this environment are held up to a very high standard. It comes as no surprise that runtime performance is the top priority in language selection when an hour of computation costs thousands of dollars.[Lud11] The focus on raw power led to C and Fortran having an almost monopolistic position in the industry, because their execution speed is nearly unmatched.

However programming in these rather antique languages can be rather difficult. Although they are still in active development, their long lifespans resulted in sometimes unintuitive syntax and large amounts of historical debt. Especially C's *undefined behaviour* often causes inexperienced programmers to write unreliable code which is unnecessarily dependant on implementation details of a specific compiler. Understanding and maintaining these programs requires deep knowledge of memory layout and other technical details.

Considering the fact that scientific applications are often written by scientist without a concrete background in computer science it is evident that the current situation is less than ideal. There have been various efforts to make programming languages more accessible in the recent years but unfortunately none of the newly emerged ones have been successful in establishing themselves in the HPC community to this day. Although many features and concepts have found their way in newer revision of C and Fortran standards most of them feel tacked (//wording) on and are not well integrated into the core languages.

One example for this is the common practice of testing. Specifically with the growing popularity of test driven development it became vital to the development process to be able to quickly and regularly execute a set of tests to verify growing implementations as they are developed. Of course there are also testing frameworks and libraries for Fortran and C but since these languages lack deep integration of testing concepts, they often

require a lot of setup and boilerplate code and are generally not that pleasant to work with. In contrast for example the Go programming language includes a complete testing framework with the ability to perform benchmarks, execute global setup/teardown work and even do basic output testing. [mai] (// cite vs footnote) Maybe most important all this is available via a single executable go test which may be easily integrated in scripts or other parts of the workflow.

While testing is just one example there are a lot of “best practices” and techniques which can greatly increase both developer productivity and code quality but require a language-level integration to work best. Combined with the advancements in type system theory and compiler construction both C and Fortran’s feature sets look very dated. With this in mind it is time to evaluate new potential successors of the two giants of high-performance computing.

1.2. Goals of this thesis

This thesis aims to evaluate Rust and Go as potential programming languages in the high-performance computing environment. The comparison is based on a port of an existing parallel application originally written in Python. It was rewritten in the two language candidates as well as C to have a fair base to compare to. Since libraries for interprocess communication in Rust and Go are nowhere near production-ready this thesis will focus on shared memory parallelization to avoid unfair bias based solely on the quality of the supporting library ecosystem.

The final application is a simplified version of the original but will behave nearly identical. It will use similar input formats and write result to a reduced custom output format for later analysis. To reduce complexity it does not support additional commandline arguments and has limited error handling regarding in- and output.

While performance will be the main concern additional software metrics will also be reviewed to measure the complexity and overall quality of the produced applications. Another aspect to review is the tool support and ease of development.

2. State of the art

This chapter describes the current state of the art in high-performance computing. The dominance of Fortran and C is explained and questioned in section 2.1. After that all considered language candidates are introduced and characterized.

- state of C and Fortran (section name?)
- technological advancements in low level languages - static analysis - .. -> But no real adaption possible, because language level support is missing (already included in introduction)

2.1. Programming paradigms in Fortran and C

As stated in section 1.1 high-performance computing is largely dominated by C and Fortran and although their trademark is mostly performance these two languages achieve this in very different ways.

To understand why a new language is needed it is essential to understand the shortcomings of these programming veterans. (//language?)

// remove or rephrase this (maybe as the last paragraph of this section?) The main drawback of both languages is their age. Although new revisions are regularly aproves Fortran and C strive to be backwards compatible for the most part. This has some very serious consequences especially in their respective syntax. A lot of features of newer standards are integrated suboptimally to preserve backwards compatability.

// Candidates here for now might need another chapter for those

2.2. Language candidates

As previously stated Go and Rust were chosen to be evaluated in the context of high-performance computing. This section aims to provide a rough overview of all possible language candidates that were considered for further evaluation in this thesis.

Python

Python is an interpreted general-purpose programming language which aims to be very expressive and flexible. Compared with C and Fortran which sacrifice feature richness for performance, Python's huge standard library combined with the automatic memory management offers a low border of entry and quick prototyping capabilities.

As a matter of fact many introductory computer science courses at universities in the United States recently switched from Java to Python as their first programming language. [Guo14; Lub14] This allows the students to focus on core concepts of coding and algorithms instead of distracting boilerplate code.

```
1 # Function signatures consist only of one keyword (def)
2 def fizzbuzz(start, end):
3     # Nested function definition
4     def fizzbuff_from_int(i):
5         entry = ''
6         if i%3 == 0:
7             entry += "Fizz"
8         if i%5 == 0:
9             entry += "Buzz"
10        # empty string evaluates to false (useable in
11        ↪ conditions)
12        if not entry
13            entry = str(i)
14        return entry
15    # List comprehensions are the pythonic way of composing
16    ↪ lists
17    return [int_to_fizzbuzz(i) for i in range(start, end+1)]
```

Listing 2.1: FizzBuzz in Python 3.4

In addition to the very extensive standard library the Python community has created a lot of open source projects aiming to support especially scientific applications. There is NumPy¹ which offers efficient implementations for multidimensional arrays and common numeric algorithms like Fourier transforms or MPI4Py², an MPI abstraction layer able to interface with various backends like OpenMPI or MPICH. Especially the existence of the latter shows the ongoing attempts to use Python in a cluster environment and there have been successful examples of scientific high-performance applications using these libraries(//need ref).

Unfortunately dynamic typing and automatic memory management come at a rather high price. The speed of raw numeric algorithms written in plain Python is almost always

¹ <http://www.numpy.org>

² <http://www.mpi4py.scipy.org>

orders of magnitude slower than implementations in C or Fortran. As a consequence nearly all of the mentioned libraries implement the critical routines in C and focus in optimizing the interop (// wording) experience between the two languages. This often means one needs to make tradeoffs between idiomatic Python - which might not be transferable to the foreign language - and maximum performance. As a result performance critical Python code often looks like it's equivalent written in a statically typed language. The more terseness Python loses because of this, the less desirable it becomes to use in high-performance computing since one could just fall back to C for a similar experience.

In conclusion Python was not chosen to be further evaluated because of the mentioned lack of performance (in pure Python). This might change with some new implementations emerging recently though. Most of the problems discussed here are present in all stable Python implementations today (most notably *Cython* and *PyPy*) but new projects aim to improve the execution speed in various ways. *Medusa* compiles Python code to Google's Dart to make use of the underlying virtual machine. Although these ventures are still in early phases of development, first early benchmarks promise drastic performance improvements. Once Python can achieve similar execution speed to native code it will become a serious competitor in the high performance area.

Erlang

Erlang is a relatively niche programming language originally designed for the use in telephony applications. It features a high focus on concurrency and a garbage collector which is enabled through the execution inside the BEAM virtual machine. Today it is most often used in soft real-time computing ³ because of its error tolerance, hot code reload capabilities and lock-free concurrency support. [CT09]

Erlang has a very unique and specialized syntax which is very different from C-like languages. It abstains from using any kind of parentheses as block delimiters and instead uses a mix of periods, semicolons, commas and arrows (`->`). Unfortunately the rules for applying these symbols are not very intuitive and may even seem random for newcomers at times.

One core concept of Erlang is the idea of processes. These lightweight primitives of the language are provided by the virtual machine are neither direct mappings of operating system threads nor processes. On the one hand they are cheap to create and destruct (like threads) but do not share any address space or other state (like processes). Because of this the only way to communicate is through message passing which can be handled via the `receive` keyword and send via the `!` operator. [Arm03; CT09]

```
1 %% Module example (this must match the filename - '.erl')
2 -module(example).
3 %% This module exports two functions: start and codeswitch
```

³ see https://en.wikipedia.org/wiki/Real-time_computing

```

4 %% The number after each function represents the param count
5 -export([start/0, codeswitch/1]).
6
7 start() -> loop(0).
8
9 loop(Sum) ->
10     % Match on first received message in process mailbox
11     receive
12         {increment, Count} ->
13             loop(Sum+Count);
14         {counter, Pid} ->
15             % Send current value of Sum to PID
16             Pid ! {counter, Sum},
17             loop(Sum);
18         code_switch ->
19             % Explicitly use the latest version of the function
20             % => hot code reload
21             ?MODULE:codeswitch(Sum)
22     end.
23
24 codeswitch(Sum) -> loop(Sum).

```

Listing 2.2: Erlang example

Conceptionally Erlang offers various constructs known from functional languages like pattern matching, clause based function definition and immutable variables but the language as a whole is not purely functional. Rather each erlang process in itself (ideally) behaves pure (meaning the result of a function depends solely on its input) while the collection of processes interacting with each other through messages of course contain state and side effects.

Erlang was considered as a possible candidate for HPC because of its concurrency capabilities. The fact that processes are a core part of the language and are rather cheap in both creation and destruction seems ideal for high performance applications often demanding enormous amounts of parallelism. Sadly Erlang suffers from what one might call over specialization. The well adapted type system makes it very suited for tasks where concurrency is essential like serverside request management, task scheduling and other services with high connection fluctuation, but “The ease of concurrency doesn’t make up for the difficulty in interfacing with other languages.”[Dow11] Even advocates of Erlang say they would not use Erlang for regular business logic. In high performance computing most of the processing time is spent solving numeric problems. These are of course parallelized to increase effectiveness but the concurrency aspect is often not really inherent to the problem itself. Because of this Erlang’s concurrency capabilities just do not outweigh its numeric slowness for traditional HPC problems.[Héb13]

Go

Go is a relatively young programming language which focusses on simplicity and clarity while not sacrificing too much performance. Initially developed by Google it aims to “make it easy to build simple, reliable and efficient software” (//cite). It is statically typed, offers a garbage collector, basic type inference and a large standard library. Go’s syntax is loosely inspired by C but made some major changes like removing the mandatory semicolon at the end of commands and changing the order of types and identifiers. It was chosen as a candidate because it provides simple concurrency primitives as part of the language (so called *goroutines*) while having a familiar syntax and being reasonably performant [Dox12]. It also compiles to native code without external dependencies which makes it usable on cluster computers without many additional libraries installed.

The chosen code example demonstrates two key features which are essential to concurrent programming in Go - the already mentioned goroutines as well as channels used for synchronization purposes. These provide a way to communicate with running goroutines via message passing.

```
1 package main
2
3 import "fmt"
4
5 // Number of goroutines to start
6 const GOROUTINES = 4
7
8 func helloWorldConcurrent() {
9     // Create a channel to track completion
10    c := make(chan int)
11
12    for i := 0; i < GOROUTINES; i++ {
13        // Start a goroutine
14        go func(nr int) {
15            fmt.Printf("Hello from routine %v", nr)
16            // Signalize completion via channel
17            c <- 1
18        }(i)
19    }
20
21    for i := 0; i < GOROUTINES; i++ {
22        // Wait for completion of all goroutines
23        <-c
24    }
25 }
```

Listing 2.3: Go concurrency example

Initially developed for server scenarios Go has seen production use in many different areas. At Google it is used for various internal project such as the download service “dl.google.com” which has been completely rewritten from C++ to Go in 2012. The new version can handle more bandwidth while using less memory. It is also notable that the Go codebase is about half the size of the legacy application with increased test coverage and performance [Fit13].

While Go’s focus on simplicity is admirable it has also been its greatest point of criticism. The language feature set is very carefully selected and rarely extended. It even misses some of the most natural constructs which a programmer might expect in a reasonably high-level language - the main example for this being generics. As of the time of this writing Go does not offer the common concept of generic structs (or classes) and the authors have stated this is not a big priority at the moment.

One other important fact - especially for high-performance computing - is the mandatory garbage collector. Go completely takes the burden of memory management out of the hands of the programmer and relies on the embedded runtime to efficiently perform this job. This makes it impossible to predictably allocate and release memory which can lead to performance loss. This also means the Go runtime has to be statically linked into every application. Although that might not be important for bigger codebases it increases the binary size considerably.

In the end Go was mainly chosen to be evaluated further because of promised “simple” parallelism via goroutines. It will probably not directly compete with C in execution performance but the great toolchain and simplified concurrency might outshine the performance loss.

Rust

The last candidate discussed in this chapter is Rust. Developed in the open but strongly backed by Mozilla Rust aims to directly compete with C and C++ as a systems language. It focuses on memory safety which is checked and verified at compile without (or with minimal) impact on runtime performance. Rust compiles to native code using a custom fork of the popular LLVM⁴ as backend and is compatible to common tools like **The GNU Project Debugger** (*gdb*)⁵ which makes integration into existing workflows a bit easier.

Out of the here discussed languages Rust is closest to C while attempting to fix common mistakes made possible by its loose standard allowing undefined behaviour. (//wording?) Memory safety is enforced through a very sophisticated model of ownership. It is based on common concepts which are already employed on concurrent applications but integrates them on a language level and enforces them at compile time. The basic rule is that every resource in an applications (for example allocated memory or file handles) has exactly

⁴ <http://www.llvm.org>

⁵ <http://www.gnu.org/software/gdb/>

one *owner* at a time. To share access to a resource one can use references denoted by a `&`. These can be seen as pointers in C with the additional caveat that they are readonly. To gain mutable access to a resource one must acquire a mutable reference via `&mut`. To ensure memory safety a special part of the compiler, the *borrow checker*, validates that there is never more than one mutable reference to the same resource. This effectively prevents mutable aliasing which in turn rules out a whole class of errors like iterator invalidation. It is important to remember that these checks are a “zero cost abstraction” which means they do not have any runtime overhead but enforce additional security at compile time through static analysis.

Another core aspect of Rust are *lifetimes*. As many other programming languages Rust has scopes introduced by blocks for example function and loop bodies or arbitrary scopes opened and closed by curly braces. Combined with the ownership system the compiler can exactly determine when the owner of a resource gets out of scope and call the appropriate destructor (called `drop` in Rust). This technique is called “Resource acquisition is initialization” [Str94, p. 389]. Unlike in C++ it is not limited to stack allocated objects since the compiler can rely on the ownership system to verify that no references to a resource are left when its owner gets out of scope. It is therefore safe to drop.

```
1 // Immutability per default, Option type built-in -> no null
2 fn example(val: &i32, mutable: &mut i32) -> Option<String> {
3     // Pattern matching
4     match *val {
5         /* Ranges types (x ... y notation),
6          * Powerful macro system (called via 'macro!()') */
7         v @ 1 ... 5 => Some(format!("In [1, 5]: {}", v)),
8         // Conditional matching
9         v if v < 10 => Some(format!("In [6,10): {}", v)),
10        // Constant matching
11        10          => Some("Exactly 10".to_string()),
12        /* Exhaustiveness checks at compile time,
13         * '_' matches everything */
14        -           => None
15    }
16    // statements are expressions -> no need for 'return'
17 }
```

Listing 2.4: Rust example

Although Rust focusses on performance and safety it also adopted some functional concepts like *pattern matching* and the `Option` type. Combined with `range` expressions and macros which operate on syntax level coding in Rust often feels like in a scripting language which is just very performant. This was also the main reason it was chosen to be further evaluated. Rust focusses on safety while not sacrificing any performance in the

process. Most of the checks happen at compile time making the resulting binary often close or on par with equivalent C programs. It also has the advantage of being still in development⁶ so concepts which did not work out can be quickly changed or completely dropped. But Rust's immaturity is also its greatest weakness. The

Comparison

Rust	Python	Erlang	Go
Execution model	interpreted	compiled to bytecode	compiled to native
compiled to native code			
Advantages	adv go	adv rust	
Disadvantages	speed	obscure syntax	mandatory runtime
disadv rust			
Relative speed	slow to average	average to fast	speed go
speed rust			

2.3. Related work

The search for new programming languages which are fit for HPC is not a recently developing trend. There have been multiple studies and evaluations but so far none of the proposed languages have gained enough traction to receive widespread adoption. Also most reports focused on the execution performance without really considering additional software metrics or developer productivity. [Nan+13] at least adds lines of code and development time to the equation but both of these metrics only allow for superficial conclusions about the code quality.

From the here presented candidates Go in particular has been compared to other approaches to parallel programming with mixed results. Although its regular execution speed is somewhat lacking [Mit14] showed the highest speedup from parallelization amongst the evaluated languages.

// needs more work

// really mention this here without ? Rust on the other hand has not been seriously evaluated in the HPC environment.

⁶ current version being 1.0.0-ALPHA-2 at the time of this writing

3. Approach

The first section of the third chapter describes the existing application the evaluation is based on. In addition the implementation process is illustrated and methods the comparison of languages is based on are introduced.

3.1. Overview: Streets4MPI

As stated in section 1.2 the language evaluation is based on a reimplementation of an existing parallel program originally written in Python. The application in question is *streets4MPI* - a traffic simulation using OpenStreetMap files as input and calculating shortest routes via Dijkstra's algorithm. Streets4MPI was implemented in scope of the module "Parallel Programming Project" in Spring 2012. It was written by Julian Fietkau and Joachim Nitschke and makes heavy use of the various libraries of the Python ecosystem.

//TODO: more description

The original project contained the main simulation as well as a visualization script. [FN12, p. 3] For the purpose of this thesis the visualization was omitted and the evaluation is only based on software quality metrics as well as the result's correctness. //TODO: really cite page?

3.2. Differences to the base implementation

Although the evaluated implementations are based on the original Streets4MPI, there are some key differences (mainly to reduce complexity). This section gives a brief overview over the most important aspects that have been changed and describes both the original application as well as the the redeveloped version.

In the remaining part of the thesis the different implementations will be referenced quite frequently. For brevity each application will be called by its name "streets4<language>" for example "streets4go". //wording..

Input format

The original implementation used the somewhat dated OpenStreetMap XML format ¹ as input which is parsed by *imposm.parser* ². Streets4MPI then builds a directed graph via the *python-graph* ³ library to base the simulation on.[FN12]

The reimplementations require the input to be in “.osm.pbf” format. This newer version of the OpenStreetMap format is based on Google’s *Protocol Buffers* ⁴ and is superior to the XML variant in both size and speed [Pro]. It also simplifies multilanguage development because the code performing the actual parsing is auto generated from a language independent description file and there are protobuf backends for C, Rust and Go which can perform that generation.

Simulation

The simulation in streets4MPI is based on randomly generated trips in the directed graph that was build from the input. For these trips the shortest path is calculated by Dijkstra’s algorithm (/cite). To avoid oscillation a random factor called “jam tolerance” is introduced. Then after some time has passed in the simulation existing streets get expanded or shut down depending on the usage. The reimplementations to compare also perform trip based simulation but to without the added randomness and street modification. Also a variant of Diskstra’s algorithm was implemented in all three languages in a benchmarkable way so it can be compared separately. The concept is based on the DIJKSTRA-NODEC algorithm as seen in [Che+07, p. 16] This is done to rule out possible errors in the complete streets4x implementations and be able to directly compare the algorithmic part of the application.

Concurrency

The base application parallelizes its calculations on multiple processes that communicate via message passing. This is achieved with the aforementioned MPI4Py library which delegates to a native MPI implementation installed on the system. If no supported implementation is found it falls back to a pure Python solution although the native one should be preferred for maximum performance.

Although Rust as well as Go can integrate decently with existing native code, the reimplementations will be limited to shared memory parallelization on threads. This was mostly decided to evaluate and compare the language inherent concurrency constructs rather than the quality of the forrign funtion interface. To achieve fair comparability

¹ http://wiki.openstreetmap.org/wiki/OSM_XML

² <http://imposm.org/docs/imposm.parser/latest/>

³ <https://code.google.com/p/python-graph/>

⁴ <https://developers.google.com/protocol-buffers/>

streets4c will use OpenMP ⁵ as it is the de facto standard for simple thread parallelization in C.

3.3. Implementation process

The implementation process was performed iteratively. Certain milestones were defined and implemented in all three languages. The process only proceeded to the next phase when the previous milestone was reached in all applications. This approach was chosen to allow for a fair comparison of the different stages of development. If the implementations would have been developed one after another to completion, this might have introduced a certain bias to the evaluation because of possible knowledge about the problem acquired in a previous language translating to faster results in the next one.

For each stage various characteristics were captured and compared to highlight the languages' features and performance in the various areas

3.3.0. [Setting up the project]

The first stage of development was to create project skeletons and infrastructure for the future development. The milestone was to have a working environment in place where the sample application could be build and executed. While this is certainly not the most important or even interesting part it did show the differences in comfort between the various toolchains.

3.3.1. Counting nodes, ways and relations in an .osm.pbf file

The first real milestone was to read a .osm.pbf file and count all nodes, ways and relations in it. This was done to get familiar with the required libraries and the file format in general. The time recorded began from the initial project created in stage 0 and finished after the milestone was reached. As this is the most input and output intensive stage it already showed some key differences between the candidates both in speed as well as memory consumption.

3.3.2. Building a basic graph representation for the simulation

The next goal was to conceptionally build the graph and related structures the simulation would later operate on. This involved thinking about the relation between edges and nodes as well as the choice of various containers to store the objects efficiently while also keeping access simple. This milestone tested the language's standard libraries and ability to express functions in a compact way (lambda..)

⁵ <http://www.openmp.org>

3.3.3. Verifying the representation and algorithm

After the base structure to represent graphs and calculate shortest paths is in place it was time to validate the implementations. Unfortunately the OSM data used in the first phase 3.3.1 contained too much nodes and ways to be able to efficiently verify any computed results. Therefore a small example graph was manually populated and fed to the algorithm. However unlike the final application the start and end nodes were fixed to allow for comparable results.

4. Implementation

This chapter describes the implementation process for all three compared languages. It is divided in sections based on the development milestones defined in the previous chapter 3.3.

4.0. Project setup

As all applications written for this thesis have been developed on Linux the setup instructions are for this operating system. They **should** work on *nix as well but there is no definite guarantee this is the case. Also each section assumes the toolchains for the various languages are installed as this is largely different based on what operating system and on Linux which distribution is used. It is therefore not covered in this thesis.

4.0.1. C

//TODO: Add footnotes/links

The buildtool for streets4C is GNU *make* with a simple handcrafted **Makefile**. It was chosen to strike a balance between full blown build systems like *Autotools* or *CMake* and manual compilation. The setup steps required for this configuration are relatively straight forward and listed below.

```
$ mkdir -p streets4c
$ cd streets4c
$ vim main.c
$ vim Makefile
$ make && ./streets4c
```

Listing 4.1: Project setup: streets4C

After creatig a new directory for the application a **Makefile** and a sourcefile are created. **main.c** contains just a bare bones main method while the **Makefile** uses basic rules to compile an executable named streets4c with various optimization flags.

All in all the setup in C is quite painless although manual. The only caveat are Makefiles. They may be simple for small projects without real dependencies but as soon as different

source and object files are involved in the compilation process they can get quite confusing. At that point the mentioned build systems might prove their worth in generating the Makefile(s).

4.0.2. Go

//TODO: Add footnotes/links

For Go the choice of buildtool is nonexistent. The language provides the `go` executable which is responsible for nearly the complete development cycle. It can compile your code, install arbitrary Go packages from various sources, run tests and format code just to name the most common features.

This makes Go (the language) extremely convenient since everything you want to do is probably one commandline away. For example to get a dependency one would invoke the tool like so: `go get github.com/petar/GoLLRB/llrb`

This will download the package in source form which can then be imported in any project on that machine via its fully qualified package name.

To achieve this convenience the `go` tool requires some setup work before it can be used for the first time. Because of this this section contains two setup examples.

```
$ mkdir -p streets4go
$ cd streets4go
$ vim main.go
$ go run main.go
```

Listing 4.2: Project setup: streets4Go

Listing 4.2 describes the steps that were taken to create the streets4go project inside the thesis' repository. It is pretty similar to the C version. A directory gets created then a source file containing a `main` function is created which can be build and run with a single command.

One thing to remind here is the fact, that this code does not live inside a globally set directory called GOPATH. To be able to download packages only once `go` assumes an environment variable called GOPATH to be set to a directory which it has full control over. This directory contains all source files as well as the compiled binaries all stored through a consistent naming scheme. Normally it is assumed that all Go projects live inside their own subdirectories of the GOPATH but it is possible to avoid this at the cost of some convenience.

The project that was created through the commands of listing 4.2 for example cannot be installed to the system by running `go install` since it does not reside in the correct folder instead one has to copy the compiled binary to a directory in PATH manually.

The next listing shows a more realistic workflow for creating a new Go project from scratch without any prior setup required. It assumes one starts in the directory that should be set as GOPATH and assumes `www.github.com` as code host which in reality just determines the package name. It is also important to add the export shown in the first line to any initialization file of your shell or operating system to ensure it always being available.

```
$ export GOPATH=$(pwd)
$ mkdir -p src/github.com/<user>/<project>
$ cd src/github.com/<user>/<project>
$ vim main.go
$ go run main.go
```

Listing 4.3: Full setup for new Go projects

4.0.3. Rust

//TODO: Add footnotes/links

Similar to Go Rust provides its own build system. As mentioned in the candidate introduction Rust installs its own package manager `cargo`. It functions as build system and is also capable of creating new projects.

```
$ cargo new --bin streets4rust
$ cd streets4rust
$ cargo run
```

Listing 4.4: Project setup: streets4Rust

	C	Go	Rust
SLOC	55	55	36
Development time (min)	51:18	21:16	33:09
Execution time (sec)	1,017 (-O0)	4,846 (GOMAXPROCS=1)	27,749 (-O0)
	0,994 (-O3)	1,381 (GOMAXPROCS=8)	2,722 (-O3)
Allocation count	2,390,566	11,164,068 ¹	11,373,558
Free count	2,390,566	11,000,199 ²	11,373,557 ³

Table 4.1.: Milestone 1: Counting nodes, ways and relations

	C	Go	Rust
SLOC	385	196	170
Development time (hours)	02:30:32	01:06:06	01:14:28

Table 4.2.: Milestone 2: Building a basic graph representation

4.0.4. Comparison

4.1. Counting nodes, ways and relations in an .osm.pbf file

4.2. Building a basic graph representation for the simulation

Milestones:

- Count all nodes, ways and relations in hamburg-latest.osm.pbf
 - Rust:
 - SLOC: 36
 - dev-time: 1989 -> 33:09 min
 - run-time:
 - -O0: target/streets4rust ../osm/hamburg-latest.osm.pbf 27,61s user 0
 - -O3: target/release/streets4rust ../osm/hamburg-latest.osm.pbf 2,59s
 - allocs: total heap usage: 11,373,558 allocs, 11,373,557 frees, 2,186,107
 - counts: Found 2180418 nodes, 409424 ways and 7182 relations in ../osm/har
 - notes:
 - easy dependency management, huuge optimization gains

¹ The memory statistics for Go have not been acquired by valgrind but by `runtime.MemStats` this and the fact that Go is garbage collected explain the discrepancy in allocations and frees

² See footnote 1

³ This is due to a bug in the osmpbf library used. In safe Rust code it is impossible to leak memory

-
- Go:
 - SLOC: 55 (+ helper functions)
 - dev-time: 1276 -> 21:16 min
 - run-time:
 - GOMAXPROCS=1: ./streets4go ../osm/hamburg-latest.osm.pbf 4,85s user
 - GOMAXPROCS=8: ./streets4go ../osm/hamburg-latest.osm.pbf 9,00s user
 - allocs: total heap usage: 11,164,068 allocs, 11,000,199 frees, 1,447,543
 - counts: Found 2180418 nodes, 409424 ways and 7182 relations in ../osm/har
 - notes:
 - library parallelizeable, but singlethreaded slower
 - C:
 - SLOC: 55
 - dev-time: 3078 -> 51:18 min
 - run-time:
 - -00: ./streets4c ../osm/hamburg-latest.osm.pbf 0,95s user 0,07s syst
 - -03: ./streets4c ../osm/hamburg-latest.osm.pbf 0,92s user 0,08s syst
 - allocs: total heap usage: 2,390,566 allocs, 2,390,566 frees, 372,758,206
 - counts: Found 2180418 nodes, 409424 ways and 7182 relations in ../osm/har
 - notes:
 - no library available, linking problems, freeing problems, fastest so
 - Write basic graph structure
 - C:
 - SLOC: 385 (*.c), 136 (*.h) -> 521 total
 - dev-time: 12110 - 3078 = 9032 -> 02:30:32
 - notes:
 - glib verbosity, missing generics / typesafety
 - Rust:
 - SLOC: 170 total
 - dev-time: 6457 - 1989 = 4468 -> 01:14:28
 - notes:
 - lifetime gotchas (explicit readonly of dg.g etc.)
 - Go:
 - SLOC: 196 total (including util (priorityqueue) but excluding test)
 - dev-time: 5242 - 1276 = 3966 -> 01:06:06
 - notes:
 - testability

5. Evaluation

In diesem Kapitel ...

6. Conclusion

In diesem Kapitel ...

- Only evaluated shared memory -> Multi process implementations -> C: MPI, Rust: MPI via C FFI & opaque pointer, Go: MPI via wrapper? (less idiomatic code)

Bibliography

- [Arm03] Joe Armstrong. “Making reliable distributed systems in the presence of software errors”. dissertation. Stockholm, Sweden: The Royal Institute of Technology, Department of Microelectronics and Information Technology, Dec. 2003. URL: http://www.erlang.org/download/armstrong_thesis_2003.pdf (visited on 13.03.2015) (cit. on p. 8).
- [Che+07] Mo Chen et al. *Priority Queues and Dijkstra’s Algorithm*. Technical report TR-07-54. The University of Texas at Austin, Department of Computer Science, Oct. 12, 2007. URL: <http://www3.cs.stonybrook.edu/~rezaul/papers/TR-07-54.pdf> (cit. on p. 15).
- [Cor+09] Thomas H. Cormen et al. *Introduction to Algorithms, 3rd Edition*. 3rd. The MIT Press, July 2009. ISBN: 9780262033848.
- [CT09] Francesco Cesarini and Simon Thompson. *Erlang programming. A Concurrent Approach to Software Development*. Beijing: O’Reilly, June 2009. ISBN: 9780596518189 (cit. on p. 8).
- [Dow11] Malcolm Dowse. *Erlang and First-Person Shooters*. June 2011. URL: <http://www.erlang-factory.com/upload/presentations/395/ErlangandFirst-PersonShooters.pdf> (visited on 13.03.2015) (cit. on p. 9).
- [Dox12] Caleb Doxsey. *An Introduction to Programming in Go*. Lexington, KY: CreateSpace Independent Publishing Platform, Sept. 2012. ISBN: 9781478355823 (cit. on p. 10).
- [Fit13] Bradley Joseph Fitzpatrick. *dl.google.com: Powered by Go*. July 26, 2013. URL: <http://talks.golang.org/2013/oscon-dl.slide> (visited on 18.02.2015) (cit. on p. 11).
- [FN12] Julian Fietkau and Joachim Nitschke. *Project Report: Streets4MPI*. Hamburg, May 21, 2012. URL: http://wr.informatik.uni-hamburg.de/_media/research/labs/2012/2012-05-julian_fietkau_joachim_nitschke-streets4mpi-report.pdf (visited on 18.02.2015) (cit. on pp. 14, 15).
- [Guo14] Philip Guo. *Python is Now the Most Popular Introductory Teaching Language at Top U.S. Universities*. Oct. 7, 2014. URL: <http://cacm.acm.org/blogs/blog-cacm/176450-python-is-now-the-most-popular-introductory-teaching-language-at-top-us-universities/fulltext> (visited on 07.12.2014) (cit. on p. 7).

- [Héb13] Fred Hébert. *Learn you some Erlang for great good! a beginner's guide*. San Francisco: No Starch Press, 2013. ISBN: 9781593274351. URL: <http://learnyousomeerlang.com/content> (cit. on p. 9).
- [Lub14] Bill Lubanovic. *Introducing Python. Modern Computing in Simple Packages*. 1st ed. Beijing: O'Reilly Media, Inc., Nov. 26, 2014. ISBN: 9781449359362 (cit. on p. 7).
- [Lud11] Thomas Ludwig. *The Costs of Science in the Exascale Era*. May 31, 2011. URL: http://perso.ens-lyon.fr/laurent.lefevre/greendaysparis/slides/greendaysparis_Thomas_Ludwig.pdf (visited on 02.12.2014) (cit. on p. 4).
- [mai] The Go project maintainers. *The Go programming language - Package testing*. URL: <http://golang.org/pkg/testing/> (visited on 11.02.2015) (cit. on p. 5).
- [Mit14] Sparsh Mittal. "A Study of Successive Over-relaxation Method Parallelisation over Modern HPC Languages". In: *International Journal of High Performance Computing and Networking* 7.4 (June 2014), pp. 292–298. ISSN: 1740-0562. DOI: 10.1504/IJHPCN.2014.062731 (cit. on p. 13).
- [Nan+13] Sebastian Nanz et al. "Benchmarking Usability and Performance of Multicore Languages". In: *Empirical Software Engineering and Measurement, 2013 ACM / IEEE International Symposium on*. Oct. 2013, pp. 183–192. DOI: 10.1109/ESEM.2013.10 (cit. on p. 13).
- [Pro] The OpenStreetMap Project. *OpenStreeMap Wiki - "PBF Format"*. URL: http://wiki.openstreetmap.org/wiki/PBF_Format (visited on 18.02.2015) (cit. on p. 15).
- [Str94] Bjarne Stroustrup. *The Design and Evolution of C++*. 1st ed. Addison-Wesley Professional, Apr. 1994. ISBN: 9780201543308 (cit. on p. 12).

List of Figures

List of Tables

4.1	Milestone 1: Counting nodes, ways and relations	21
4.2	Milestone 2: Building a basic graph representation	21

List of Listings

2.1	FizzBuzz in Python 3.4	7
2.2	Erlang example	8
2.3	Go concurrency example	10
2.4	Rust example	12
4.1	Project setup: streets4C	18
4.2	Project setup: streets4Go	19
4.3	Full setup for new Go projects	20
4.4	Project setup: streets4Rust	20
A.1	Output of <code>uname -a</code>	31
A.2	Output of <code>gcc --version</code>	31
A.3	Output of <code>go version</code>	31
A.4	Output of <code>rustc --version</code>	31
A.5	Output of <code>cargo --version</code>	31

Appendices

A. System configuration

```
Linux florian-arch 3.18.6-1-ARCH #1 SMP PREEMPT Sat Feb 7  
  ↪ 08:44:05 CET 2015 x86_64 GNU/Linux
```

Listing A.1: Output of `uname -a`

```
gcc (GCC) 4.9.2 20150304 (prerelease)  
Copyright (C) 2014 Free Software Foundation, Inc.  
This is free software; see the source for copying  
  ↪ conditions.  There is NO  
warranty; not even for MERCHANTABILITY or FITNESS FOR A  
  ↪ PARTICULAR PURPOSE.
```

Listing A.2: Output of `gcc --version`

```
go version go1.4.2 linux/amd64
```

Listing A.3: Output of `go version`

```
rustc 1.0.0-nightly (30e1f9a1c 2015-03-14) (built  
  ↪ 2015-03-15)
```

Listing A.4: Output of `rustc --version`

```
cargo 0.0.1-pre-nightly (07cd618 2015-03-12) (built  
  ↪ 2015-03-14)
```

Listing A.5: Output of `cargo --version`

- System configuration -> clang/gcc version -> rustc and cargo version (possibly including commit hash!) -> go version and implementation (cgo?)

Erklärung

Ich versichere, dass ich die Arbeit selbstständig verfasst und keine anderen, als die angegebenen Hilfsmittel – insbesondere keine im Quellenverzeichnis nicht benannten Internetquellen – benutzt habe, die Arbeit vorher nicht in einem anderen Prüfungsverfahren eingereicht habe und die eingereichte schriftliche Fassung der auf dem elektronischen Speichermedium entspricht.

Optional: Ich bin mit der Einstellung der Bachelor-Arbeit in den Bestand der Bibliothek des Fachbereichs Informatik einverstanden.

Hamburg, den 26.11.2014